

GHC Runtime Linker

by Example

DNS is likely some of the most broadly deployed, yet most poorly understood software components deployed in the world today. What else is as common but less understood?



Timothy Perrett @timperrett

2:12am - 19 Apr 2019

Linkers [twitter.com/timperrett/sta...](https://twitter.com/timperrett/status/111591000000000000)



Gabriel Gonzalez @GabrielG439

1:32pm - 19 Apr 2019

Linking

+

Runtime

=

?

GHCi ← Byte Code ← Source Code
— — — ↗? Compiled Code ↙ — — —

Do we even need **both**?

```
module FibSlow where

fib :: Int -> Int
fib 0 = 1
fib 1 = 1
fib n = fib (n - 1) + fib (n - 2)
```

GHCI

```
λ> :set +S
λ> :load *haskell/FibSlow
[1 of 1] Compiling Fib    ( haskell/FibSlow.hs, interpreted )
λ> fib 35
9227465
(7.14 secs, 4,897,229,624 bytes)
λ> :load haskell/FibSlow
λ> fib 35
9227465
(0.50 secs, 2,150,044,856 bytes)
```

GHCI

```
λ> :set +s
λ> :load *haskell/FibSlow
[1 of 1] Compiling Fib    ( haskell/FibSlow.hs, interpreted )
λ> fib 35
9227465
(7.14 secs, 4,897,229,624 bytes)
λ> :load haskell/FibSlow
λ> fib 35
9227465
(0.50 secs, 2,150,044,856 bytes)
```

plan

1

Linkers

Simple in theory, complex in practice.

2

Linking in runtime

There's more than one way to do it²

² **Two** ways to be precise: dynamic and static.

3

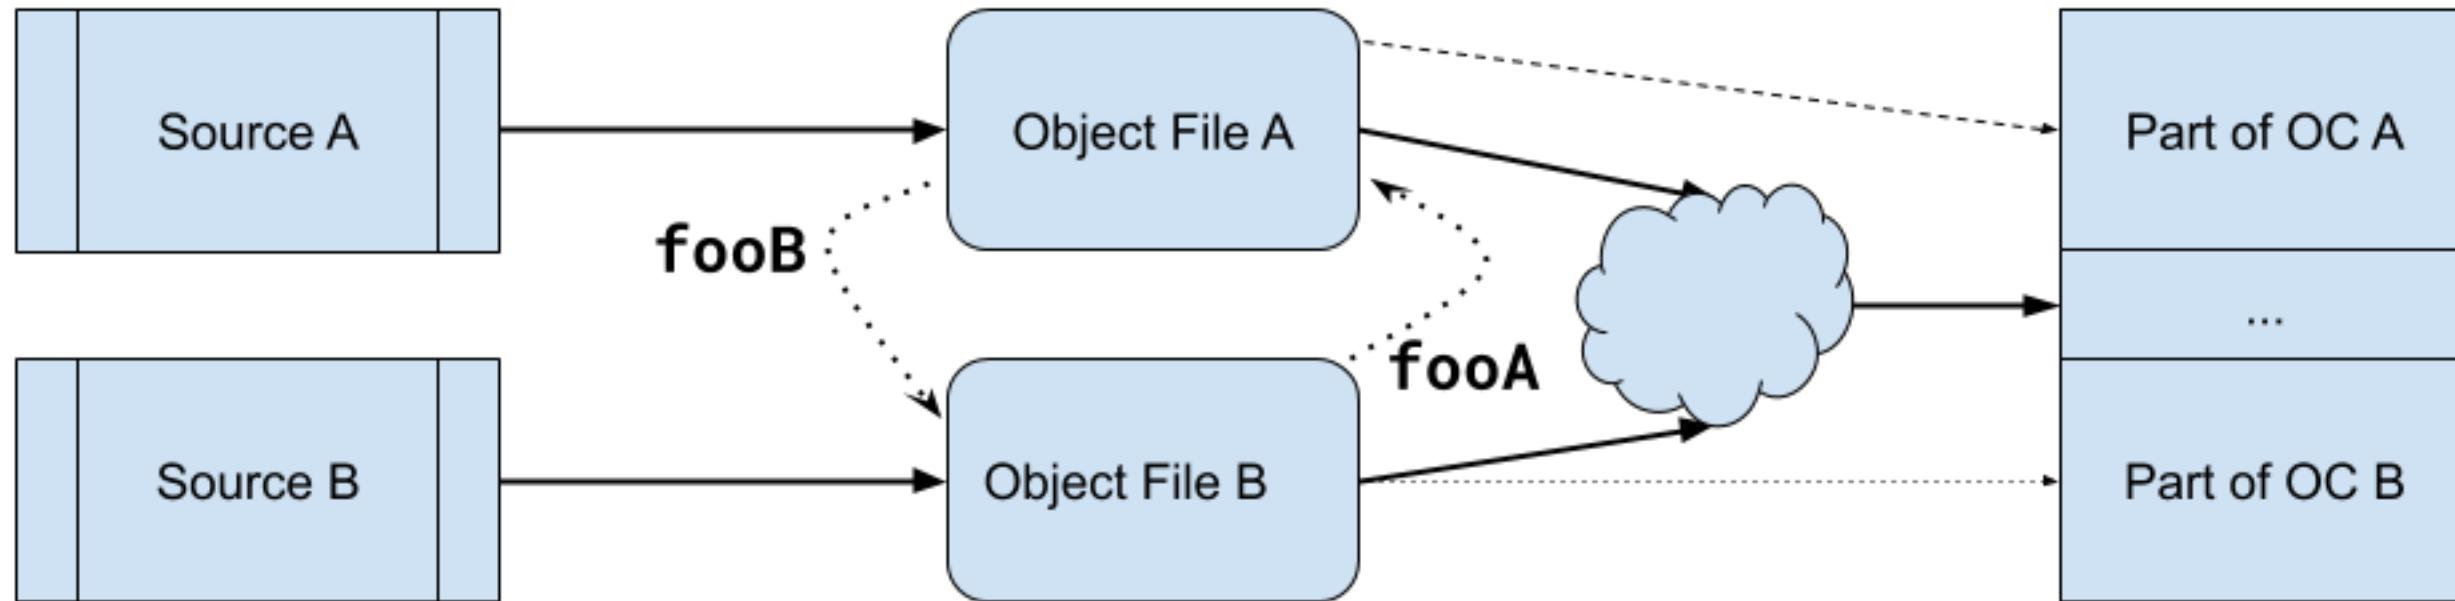


... or in other words: pros, cons, and challenges.

1

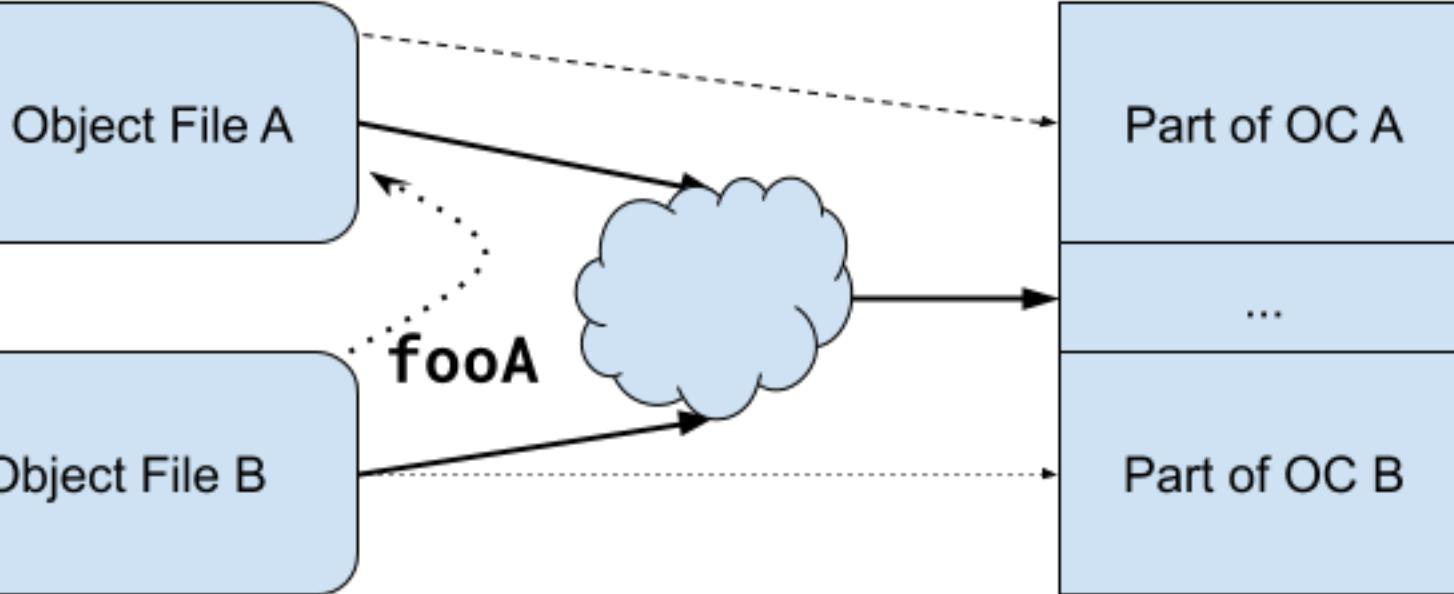
ld = linker

Compilation



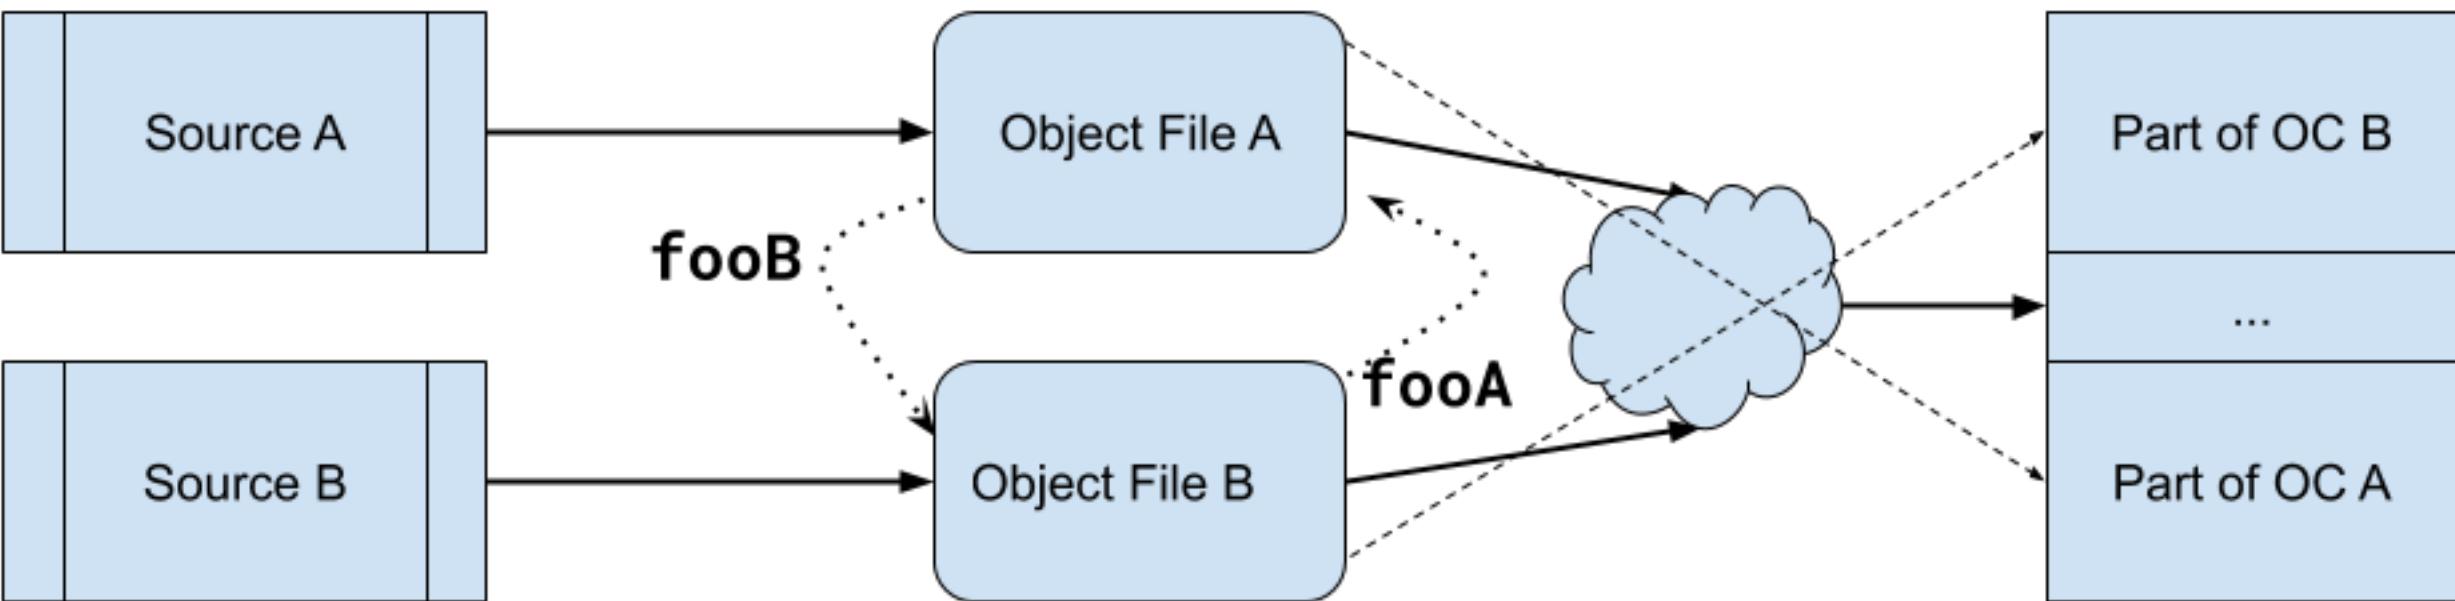
Linking

Executable



Case 1

Case 2



file:c/magic.c

```
int MAGIC = 42;
```

```
int magic(int x)
{
    return x + MAGIC;
}
```

c/magic.o:
(_TEXT,__text) section
_magic:

```
0: 55 pushq %rbp
1: 48 89 e5 movq    %rsp, %rbp
4: 03 3d 00 00 00 00 addl    _MAGIC(%rip), %edi
a: 89 f8 movl    %edi, %eax
c: 5d popq   %rbp
d: c3 retq
```

SYMBOL TABLE:

0000000000000010	g	__DATA,__data	_MAGIC
0000000000000000	g	F __TEXT,__text	_magic

Contents of (__DATA,__data) section

0000000000000010	2a 00 00 00
------------------	-------------

c/magic.o:
(__TEXT,__text) section

_magic:

0:	55	pushq	%rbp
1:	48 89 e5	movq	%rsp, %rbp
4:	03 3d 00 00 00 00	addl	<u>_MAGIC(%rip), %edi</u>
a:	89 f8	movl	%edi, %eax
c:	5d	popq	%rbp
d:	c3	retq	

SYMBOL TABLE:

0000000000000010	g	__DATA,__data	_MAGIC
0000000000000000	g	__TEXT,__text	_magic

Contents of (__DATA,__data) section

0000000000000010	2a 00 00 00
------------------	-------------

c/magic.o:
(_TEXT,__text) section
_magic:

0:	55	pushq	%rbp		
1:	48 89 e5	movq		%rsp, %rbp	
4:	03 3d 00 00 00 00	addl			_MAGIC(%rip), %edi
a:	89 f8	movl	%edi, %eax		
c:	5d	popq	%rbp		
d:	c3	retq			

SYMBOL TABLE:

0000000000000010	g	__DATA,__data	_MAGIC
0000000000000000	g	__TEXT,__text	_magic

Contents of (__DATA,__data) section

0000000000000010	2a 00 00 00
------------------	-------------

4: 03 3d 00 00 00 00 addl _MAGIC(%rip), %edi

Bytes	Value	Meaning
1	03	ADD opcode
2	3D = b00_111_101	addressing mode + register
3-6 (4)	0 stub	32bit offset from %rip

4: 03 3d 00 00 00 00 addl _MAGIC(%rip), %edi

==

4: 03 3d ?? ?? ?? ?? addl _MAGIC(%rip), %edi

```
c/magic.o:  
(__TEXT,__text) section  
_magic:  
    0: 55  pushq   %rbp  
    1: 48 89 e5      movq     %rsp, %rbp  
    4: 03 3d 00 00 00 00  addl    _MAGIC(%rip), %edi  
    a: 89 f8  movl    %edi, %eax  
    c: 5d  popq    %rbp  
    d: c3  retq
```

SYMBOL TABLE:

0000000000000010 g	__DATA,__data	_MAGIC
0000000000000000 g F	__TEXT,__text	_magic

Relocation information (__TEXT,__text) 1 entries

address	pcrel	length	extern	type	scattered	symbolnum	value
00000006	True	long	True	SIGNED	False		_MAGIC

file:c/array_magic.c

```
extern int magic(int);

void array_magic(int* arr, int len)
{
    for (int i = 0; i < len; i++) {
        arr[i] = magic(arr[i]);
    }
}
```

```

void array_magic(int* arr, int len); // rdi <- arr; esi <- len

(__TEXT,__text) section
_array_magic:
...
    a: 85 f6 testl %esi, %esi
    c: 7e 27 jle 0x35
    e: 49 89 ff movq %rdi, %r15
   11: 41 89 f6 movl %esi, %r14d
   14: 31 db xorl %ebx, %ebx
...
   20: 41 8b 3c 9f movl (%r15,%rbx,4), %edi
   24: e8 00 00 00 00 callq _magic
   29: 41 89 04 9f movl %eax, (%r15,%rbx,4)
   2d: 48 ff c3 incq %rbx
   30: 49 39 de cmpq %rbx, %r14
   33: 75 eb jne 0x20
   35: 48 83 c4 08 addq $8, %rsp

```

SYMBOL TABLE:

0000000000000000 g F __TEXT,__text _array_magic

0000000000000000 *UND* _magic

Relocation information (__TEXT,__text) 1 entries

address	pcrel	length	extern	type	scattered	symbol	num/value
---------	-------	--------	--------	------	-----------	--------	-----------

00000025	True	long	True	BRANCH	False	_magic	
----------	------	------	------	--------	-------	--------	--

```

void array_magic(int* arr, int len); // rdi <- arr; esi <- len

(__TEXT,__text) section
_array_magic:
...
a: 85 f6 testl %esi, %esi
c: 7e 27 jle 0x35
e: 49 89 ff movq %rdi, %r15
11: 41 89 f6 movl %esi, %r14d
14: 31 db xorl %ebx, %ebx
...
20: 41 8b 3c 9f movl (%r15,%rbx,4), %edi
24: e8 00 00 00 00 callq _magic
29: 41 89 04 9f movl %eax, (%r15,%rbx,4)
2d: 48 ff c3 incq %rbx
30: 49 39 de cmpq %rbx, %r14
33: 75 eb jne 0x20
35: 48 83 c4 08 addq $8, %rsp

```

SYMBOL TABLE:

0000000000000000	g	F	__TEXT,__text	_array_magic		
0000000000000000			*UND*	_magic		
Relocation information (__TEXT,__text) 1 entries						
address	pcrel	length	extern	type	scattered	symbolnum/value
00000025	True	long	True	BRANCH	False	_magic

```
c: 7e 27    jle 0x35  
e: 49 89 ff    movq    %rdi, %r15  
...  
35: 48 83 c4 08    addq    $8, %rsp  
...
```

- 7e is opcode for JLE rel8 instruction
- JLE takes a **relative offset** from the program counter
- e + 27 = 35

24: e8 00 00 00 00 callq _magic

 ^ ^-----

 | |

 32bit 0-stub

near call, 32bit displacement relative to next instr.

Relocation information (_TEXT, __text) 1 entries

address pcrel length extern type scattered symbolnum/value

00000025 True long True BRANCH False _magic

magic.o and array_magic.o combined

(__TEXT,__text) section

_magic:

```
0: 55 pushq %rbp  
1: 48 89 e5 movq %rsp, %rbp  
4: 03 3d ?? ?? ?? ?? addl _MAGIC(%rip), %edi  
a: 89 f8 movl %edi, %eax  
c: 5d popq %rbp  
d: c3 retq
```

_array_magic:

```
e: 55 pushq %rbp
```

...

```
2e: 41 8b 3c 9f movl _array_magic(%r15,%rbx,4), %edi  
32: e8 ?? ?? ?? ?? callq _magic  
37: 41 89 04 9f movl %eax, _array_magic(%r15,%rbx,4)
```

...

```
4d: c3 retq
```

Contents of (__DATA,__data) section

```
0000000000000050 2a 00 00 00
```


2

Runtime linking in GHC

-dynamic[~]

[~] dynamic way = haskell libraries and RTS are linked dynamically

```
module Foo (foo) where
```

```
foo x = x * 3
```

```
$ ghc -c Foo.hs # produces Foo.o
```

```
$ ghci
```

```
λ> :l Foo
```

```
... # Will Foo.o get used?
```

```
module Foo (foo) where
```

```
foo x = x * 3
```

```
$ ghc -c Foo.hs  
$ ghci
```

```
λ> :l Foo  
[1 of 1] Compiling Foo  
Ok, one module loaded.
```

(hs/Foo.hs, interpreted)

-dynamic by default

```
$ objdump -dylibs-used {path-to-ghc-bin}  
{path-to-ghc-bin}:  
    /usr/lib/libSystem.B.dylib  
    @rpath/libHSkeline-0.7.4.3-ghc8.6.5.dylib  
    @rpath/libHSghc-8.6.5-ghc8.6.5.dylib  
...  
    @rpath/libHSrts_thr-ghc8.6.5.dylib
```

```
$ ghc -c -dynamic Foo.hs
$ ghci
λ> :l Foo
Ok, one module loaded.
Collecting type info for 1 module(s) ...
```

```
diff -u static/Foo.s dyn/Foo.s
```

```
-static/Foo.s:  
+dyn/Foo.s:  
...  
- leaq    _stg_ap_pp_info(%rip), %rax  
+ movq    _stg_ap_pp_info(%rip), %rax  
...  
- jmp _base_GHCziNum_zt_info  
+ jmpq   *_base_GHCziNum_zt_info(%rip)  
...
```

```
diff -u <(objdump -r static/Foo.o) <(objdump -r dyn/Foo.o)
```

```
- X86_64_RELOC_BRANCH _base_GHCziNum_zt_info
- X86_64_RELOC_SIGNED _stg_ap_pp_info
...
+ X86_64_RELOC_GOT _base_GHCziNum_zt_info@GOTPCREL
+ X86_64_RELOC_GOT_LOAD _stg_ap_pp_info@GOTPCREL
```

```
$ find [ghci]  
./compiler/ghci  
./libraries/ghci
```

```
compiler/ghci
├── Linker.hs
└── LinkerTypes.hs
...
...
```

	files	blank	comment	code
Linker.hs	1	262	422	1002
Total	11	1031	1735	4703

file:Linker.hs

```
dynLinkObjs hsc_env pls objs = do
    -- Load the object files and link them
    let (objs_loaded', new_objs) = rmDupLinkables (objs_loaded pls) objs
        pls1                      = pls { objs_loaded = objs_loaded' }
        unlinkeds                  = concatMap linkableUnlinked new_objs
        wanted_objs                = map nameOfObject unlinkeds

    if interpreterDynamic (hsc_dflags hsc_env)
        then do pls2 <- dynLoadObjs hsc_env pls1 wanted_objs
                return (pls2, Succeeded)
        else do mapM_ (loadObj hsc_env) wanted_objs
                -- Link them all together
                ok <- resolveObjs hsc_env
    ...
```

file:Linker.hs

```
dynLoadObjs _           pls                      []   = return pls
dynLoadObjs hsc_env pls@PersistentLinkerState{..} objs = do
  let dflags = hsc_dflags hsc_env
  ...
  (soFile, libPath , libName) <- newTempLibName
                                dflags
                                TFL_CurrentModule (soExt platform)

  let dflags2 = dflags {..., outputFile = Just soFile }
  -- link all "loaded packages" so symbols in those can be resolved
  linkDynLib dflags2 objs pkgs_loaded

  -- if we got this far, extend the lifetime of the library file
  changeTempFilesLifetime dflags TFL_GhcSession [soFile]
  m <- loadDLL hsc_env soFile
  case m of
    Nothing -> return $! pls { temp_sos = (libPath, libName) : temp_sos }
    Just err -> panic ("Loading temp shared object failed: " ++ err)
```

file:Linker.hs

```
dynLoadObjs _           pls          []      = return pls
dynLoadObjs hsc_env pls@PersistentLinkerState{..} objs = do
  let dflags = hsc_dflags hsc_env
  ...
  (soFile, libPath , libName) <- newTempLibName
                                dflags
                                TFL_CurrentModule (soExt platform)

  let dflags2 = dflags {..., outputFile = Just soFile }
  -- link all "loaded packages" so symbols in those can be resolved
  linkDynLib dflags2 objs pkgs_loaded

  -- if we got this far, extend the lifetime of the library file
  changeTempFilesLifetime dflags TFL_GhcSession [soFile]
  m <- loadDLL hsc_env soFile
  case m of
    Nothing -> return $! pls { temp_sos = (libPath, libName) : temp_sos }
    Just err -> panic ("Loading temp shared object failed: " ++ err)
```

file:Linker.hs

```
dynLoadObjs _           pls          []      = return pls
dynLoadObjs hsc_env pls@PersistentLinkerState{..} objs = do
  let dflags = hsc_dflags hsc_env
  ...
  (soFile, libPath , libName) <- newTempLibName
                                dflags
                                TFL_CurrentModule (soExt platform)

  let dflags2 = dflags {..., outputFile = Just soFile }
  -- link all "loaded packages" so symbols in those can be resolved
  linkDynLib dflags2 objs pkgs_loaded

  -- if we got this far, extend the lifetime of the library file
  changeTempFilesLifetime dflags TFL_GhcSession [soFile]
  m <- loadDLL hsc_env soFile
  case m of
    Nothing -> return $! pls { temp_sos = (libPath, libName) : temp_sos }
    Just err -> panic ("Loading temp shared object failed: " ++ err)
```

very nice!

Where's the catch?

Example

```
-- file: Foo.hs  
module Foo (foo) where  
  
foo a = a * 3
```

```
-- file: Bar.hs  
module Bar (bar) where  
import Foo  
  
bar a = even $ foo a + 1
```

```
$ ghc --interactive +RTS -Dl  
initLinker: start  
...  
Prelude> :l Foo Bar  
Ok, two modules loaded.  
Prelude Foo> :m Foo Bar  
Prelude Foo Bar>
```

```
Prelude Foo Bar> foo 2
addDLL: dll_name = '[tmp-forlder]/libghc_1.dylib'
internal_dlopen: dll_name = '[tmp-folder]/libghc_1.dylib'
...

```

```
Prelude Foo Bar> bar 3
addDLL: dll_name = '[tmp-folder]/libghc_3.dylib'
internal_dlopen: dll_name = '[tmp-folder]/libghc_3.dylib'
...

```

```
$ PAT='.*F.*\(|Foo_foo\|Bar_bar\|).*'  
  
$ objdump -t libghc_1.dylib | grep $PAT  
l      F __TEXT,__text      _dsp_Foo_foo_info_dsp  
g      F __TEXT,__text      _Foo_foo_info  
  
$ objdump -t libghc_3.dylib | grep $PAT  
l      F __TEXT,__text      _dsp_Bar_bar_info_dsp  
g      F __TEXT,__text      _Bar_bar_info
```

```
Prelude Foo Bar> :l Bar
```

```
Ok, two modules loaded.
```

```
Prelude Bar> bar 4
```

```
addDLL: dll_name = '[tmp-folder]/libghc_5.dylib'
```

```
internal_dlopen: dll_name = '[tmp-folder]/libghc_5.dylib'
```

```
...
```

```
$ objdump -t libghc_5.dylib | grep $PAT
l      F __TEXT,__text      _dsp__Foo_foo_info_dsp
l      F __TEXT,__text      _dsp__Bar_bar_info_dsp
g      F __TEXT,__text      _Bar_bar_info
g      F __TEXT,__text      _Foo_foo_info
```

```
Prelude Bar> :l Bar
Ok, two modules loaded.
Prelude Bar> bar 5
addDLL: dll_name = '[tmp-folder]/libghc_7.dylib'
internal_dlopen: dll_name = '[tmp-folder]/libghc_7.dylib'
...

```

```
$ objdump -t libghc_7.dylib | grep $PAT
l      F __TEXT,__text      _dsp__Foo_foo_info_dsp
l      F __TEXT,__text      _dsp__Bar_bar_info_dsp
g      F __TEXT,__text      _Bar_bar_info
g      F __TEXT,__text      _Foo_foo_info
```

```
$ lsof -c GHC | grep libGHC_
```

CMD	PID	FD	TYPE	NAME
GHC	91710	txt	REG	[tmp-folder]/libGHC_1.dylib
GHC	91710	txt	REG	[tmp-folder]/libGHC_3.dylib
GHC	91710	txt	REG	[tmp-folder]/libGHC_5.dylib
GHC	91710	txt	REG	[tmp-folder]/libGHC_7.dylib



- Takes time to find what it needs.
- Never gives back what it once received.

-vanilla*

* *vanilla way* = haskell libraries and RTS are linked statically

file:Linker.hs

```
dynLinkObjs hsc_env pls objs = do
    -- Load the object files and link them
    let (objs_loaded', new_objs) = rmDupLinkables (objs_loaded pls) objs
        pls1 = pls { objs_loaded = objs_loaded' }
        unlinkededs = concatMap linkableUnlinked new_objs
        wanted_objs = map nameOfObject unlinkededs

    if interpreterDynamic (hsc_dflags hsc_env)
        then do pls2 <- dynLoadObjs hsc_env pls1 wanted_objs
                return (pls2, Succeeded)
        else do mapM_ (loadObj hsc_env) wanted_objs
                -- Link them all together
                ok <- resolveObjs hsc_env
    ...
    ...
```

```
libraries/ghci/GHCi  
└── ObjLink.hs  
...
```

	files	blank	comment	code
ObjLink	1	27	45	123
Total	12	292	414	1437

```
rts
└── Linker.c
└── LinkerInternals.h
└── linker
    ...
    └── ELFRelocs
        ...
        ├── Elf.c
        ├── Elf.h
        ├── ElfTypes.h
        ├── LoadArchive.c
        ├── M32Alloc.c
        ├── M32Alloc.h
        ├── Mach0.c
        ├── Mach0.h
        ├── Mach0Types.h
        ├── PEi386.c
        ├── PEi386.h
        ├── PEi386Types.h
        ├── SymbolExtras.c
        ├── SymbolExtras.h
        └── util.h
```

```
└── elf_compat.h
    ├── elf_got.c
    ├── elf_got.h
    ├── elf_plt.c
    ├── elf_plt.h
    ├── elf_plt_aarch64.c
    ├── elf_plt_aarch64.h
    ├── elf_plt_arm.c
    ├── elf_plt_arm.h
    ├── elf_reloc.c
    ├── elf_reloc.h
    ├── elf_reloc_aarch64.c
    ├── elf_reloc_aarch64.h
    ├── elf_util.c
    └── elf_util.h
```

A **fair** amount of **C** code!

Language files	blank	comment	code
C	15	1247	1977
C Header	19	217	265

Linker C interface

initLinker_

addLibrarySearchPath

removeLibrarySearchPath

findSystemLibrary

addDLL

loadArchive

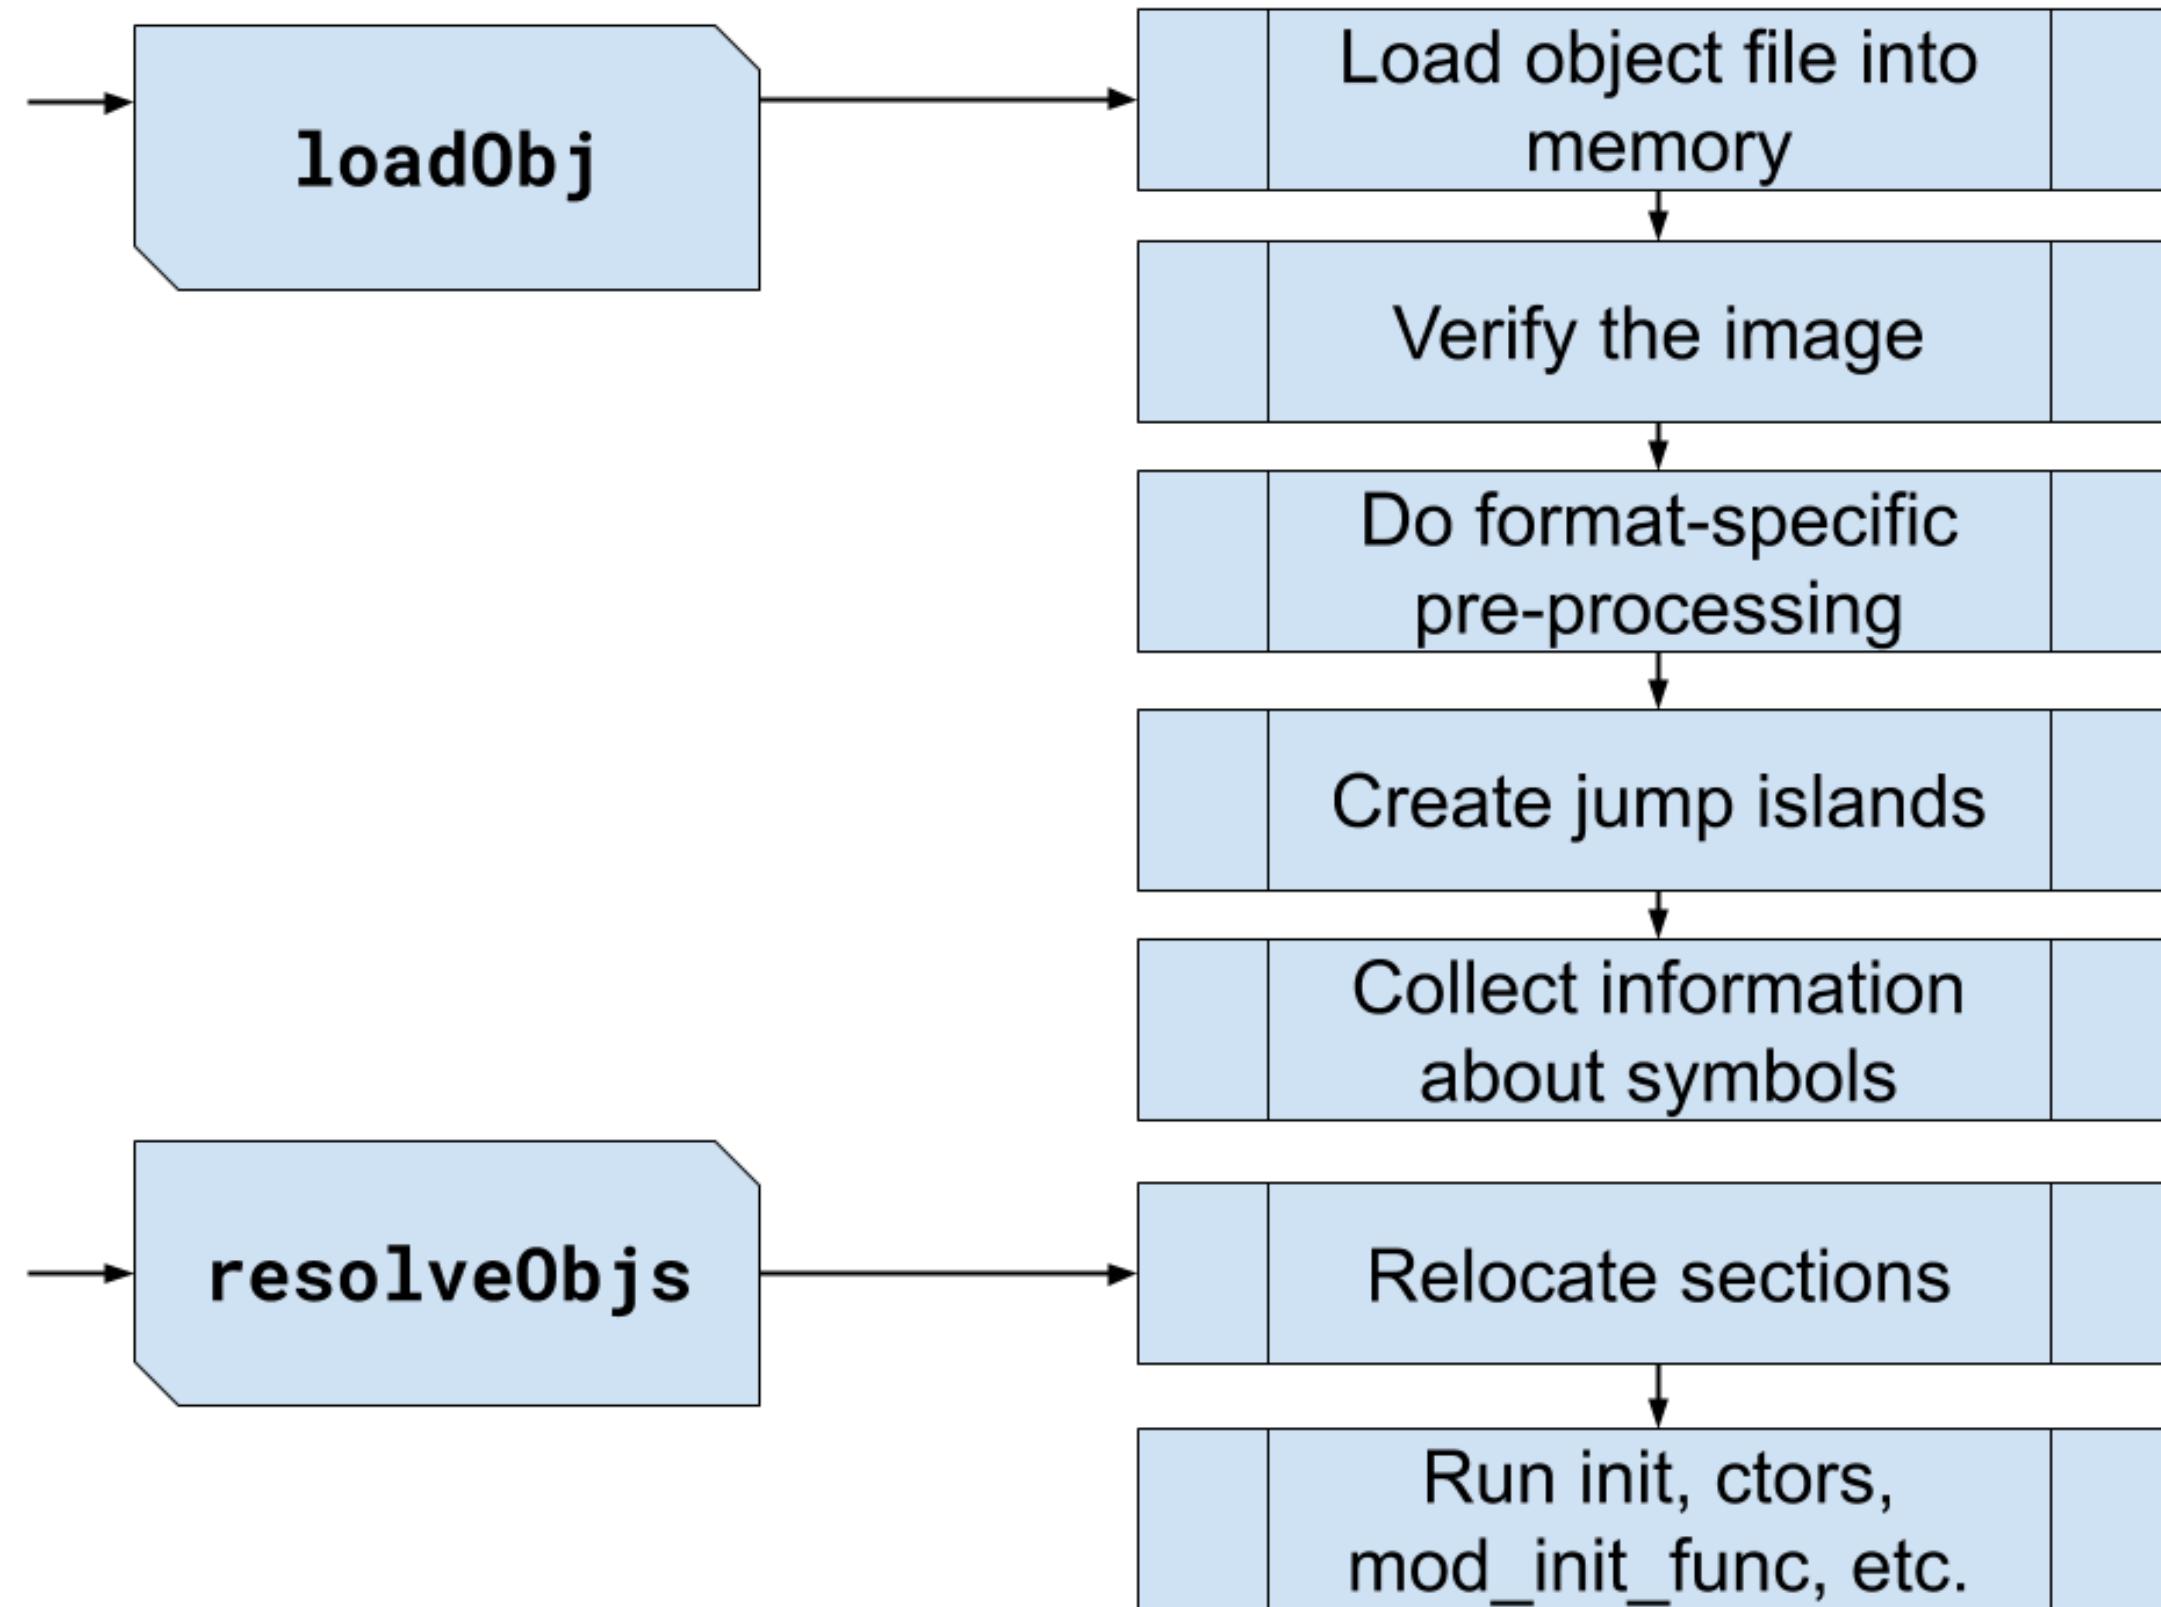
loadObj

purgeObj

unloadObj

resolveObjs

lookupSymbol



**What is the
problem? =**

⁼ ... apart from needing to deal with 10 kLOC of involved C that is hard to get right.

BUGS

A cartoon illustration of a red and black spotted dragon-like creature with large white eyes and a long tail, surrounded by flames.

Example 1

Debugging a rare^r segfault

^r About 0.02% chance of a segfault.

```
int
ocGetNames_MachO(ObjectCode* oc) // oc was mmaped earlier
{
    // within for-loop over object code (oc) sections
    section_64 *section = &oc->info->macho_sections[i];
    // for section of type ZEROFILL
    zeroFillArea = mmap(section->size, MAP_ANONYMOUS, -1, 0);
    // check mmap
    section->offset = zeroFillArea - oc->image;
    //...
}
```

```
int
ocGetNames_MachO(ObjectCode* oc) // oc was mmapped earlier
{
    // within for-loop over object code (oc) sections
    section_64 *section = &oc->info->macho_sections[i];
    // for section of type ZEROFILL
    zeroFillArea = mmap(section->size, MAP_ANONYMOUS, -1, 0);
    section->offset = zeroFillArea - oc->image;
    //...
}

// MacOS SDK/mach-o/loader.h
struct section_64 {
    uint32_t      offset;      /* file offset of this section */
};
```

What if...

the kernel allocates
a page before
`oc->image?`

? unsigned integer underflow and segfault down the road

Example 2

**Not all numbers
are created equal**

```
int
ocBuildSegments_Macho(ObjectCode *oc)
{
    size_t size_xxSegment = 0; //xx = rx, rw, zerofill

    for (int i = 0; i < oc->n_sections; i++) {
        Mach0Section *macho = &oc->info->macho_sections[i];
        size_t alignment = 1 << macho->align;

        size_xxSegment = roundUpToAlign(size_xxSegment, alignment);
        size_xxSegment += macho->size;
    }

    size_compound = roundUpToPage(size_xxSegment) + ...;
    mem = mmap(size_compound, MAP_ANON, -1, 0);
    //...
}
```

```
int
ocBuildSegments_Macho(ObjectCode *oc)
{
    size_t size_xxSegment = 0; //xx = rx, rw, zerofill

    for (int i = 0; i < oc->n_sections; i++) {
        MachOSection *macho = &oc->info->macho_sections[i];
        size_t alignment = 1 << macho->align;

        size_xxSegment = roundUpToAlign(size_xxSegment, alignment);
        size_xxSegment += macho->size;
    }

    size_compound = roundUpToPage(size_xxSegment) + ...;
    mem = mmap(size_compound, MAP_ANON, -1, 0);
    //...
}
```

What if...

**all sections in
the object file
are empty??**

?? GHC-16701 -- mmap 0 bytes at 0x0

Bigger issue

GHC-13624!

`loadobj()` does not respect alignment

! Not fixed yet 😭

Unlucky SSE instruction

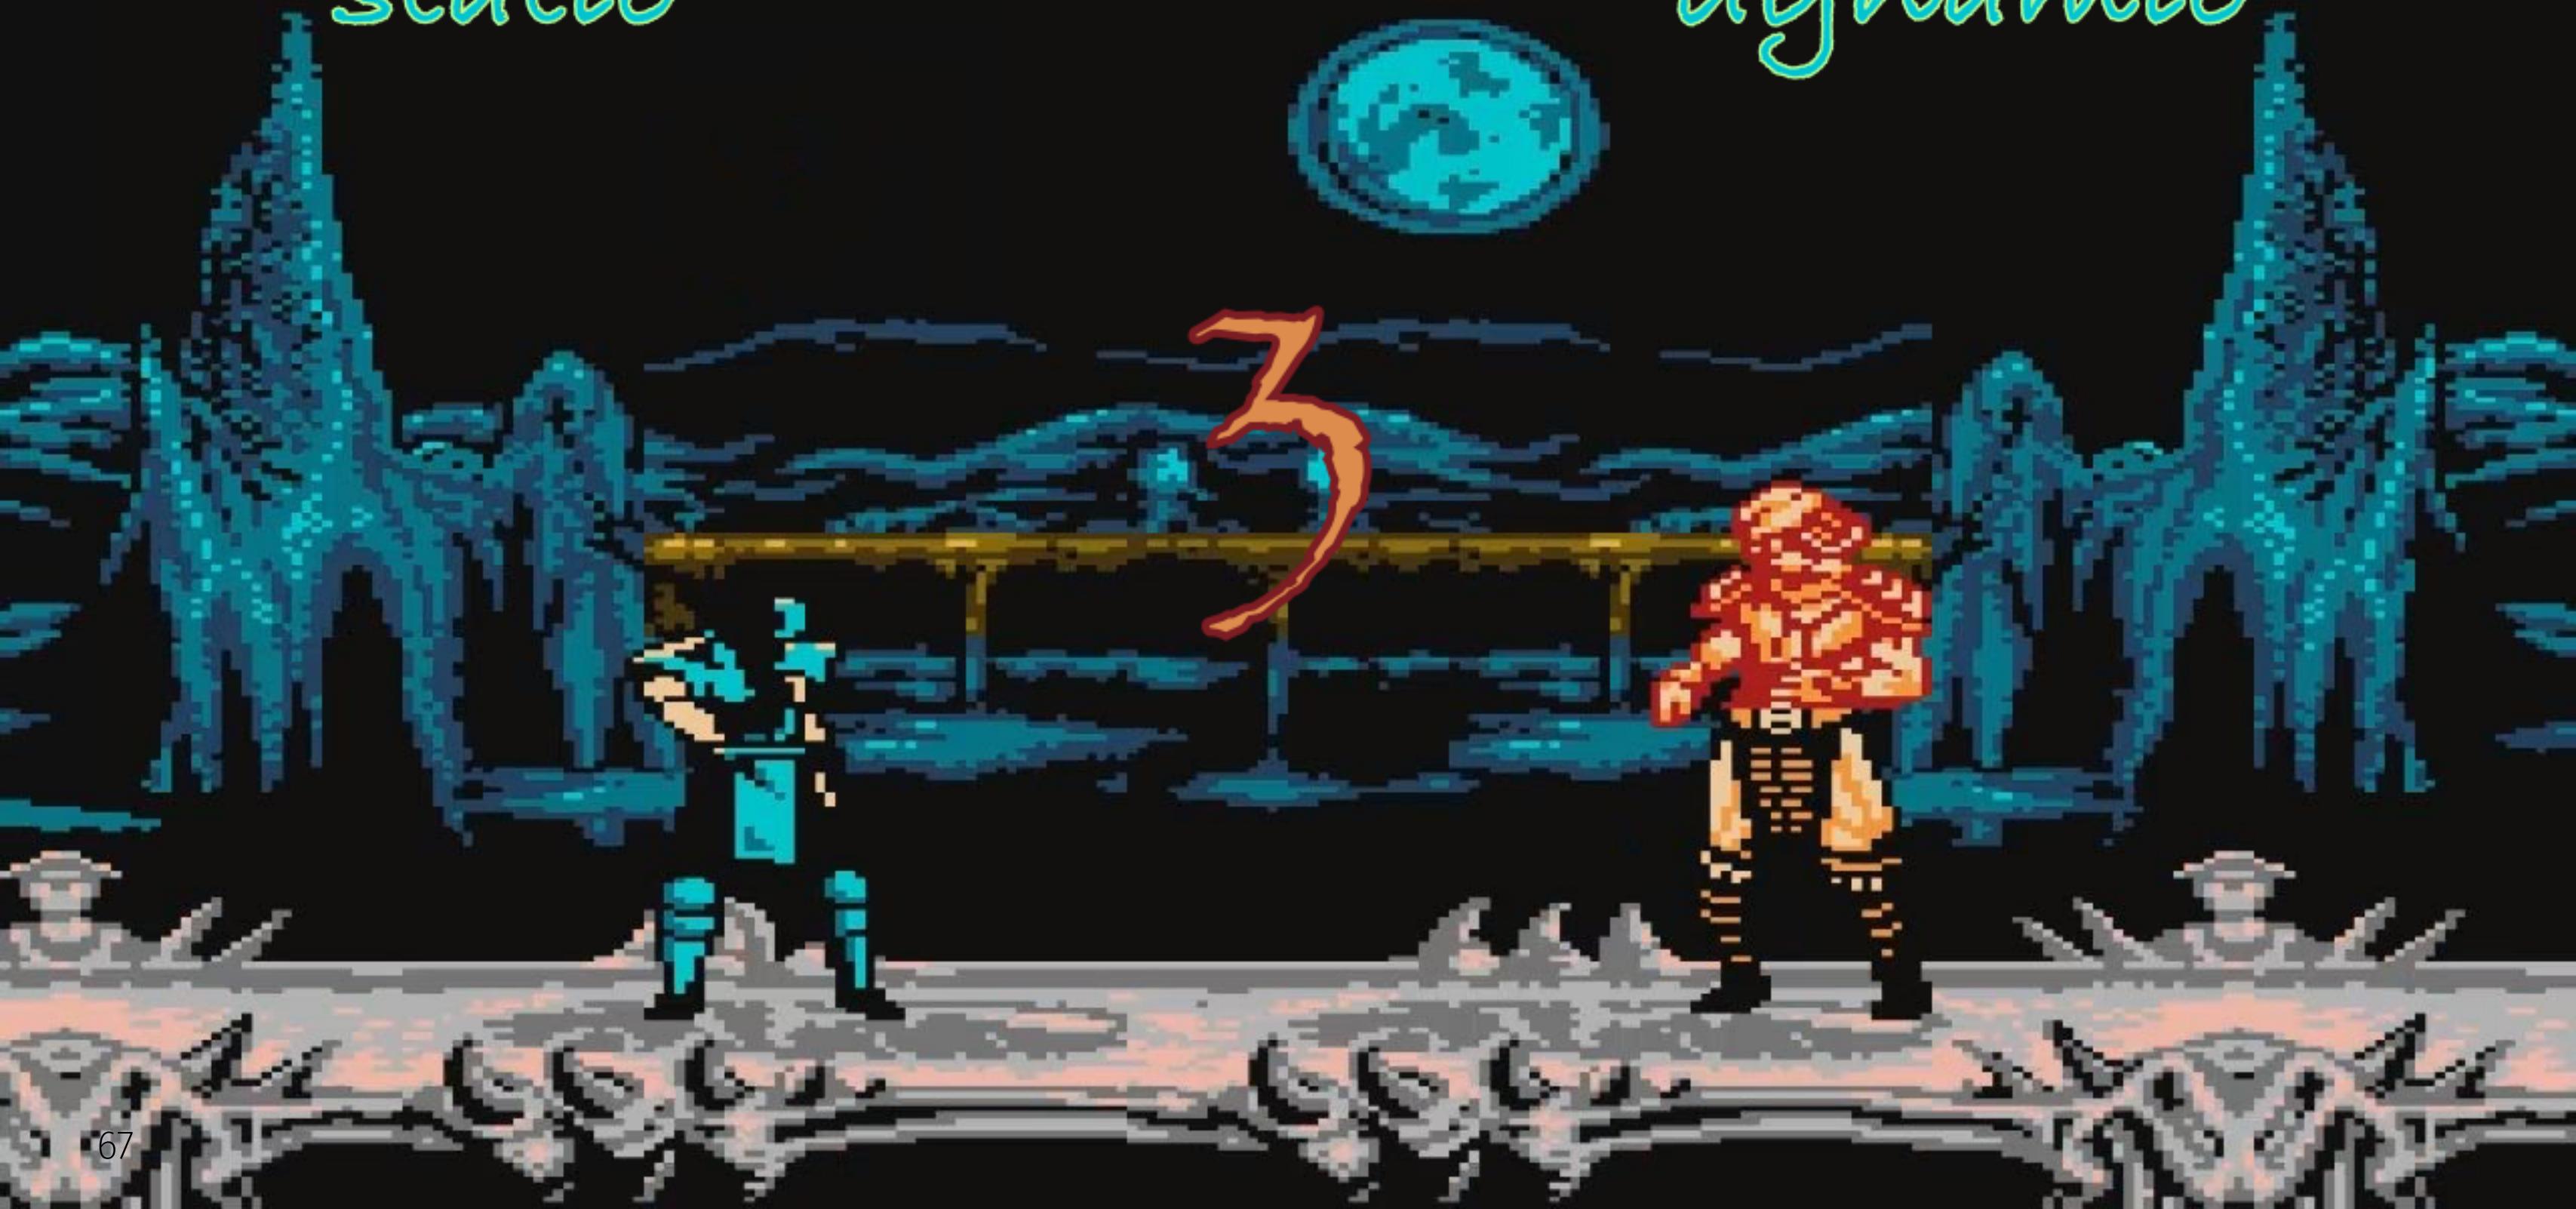
=

SIGBUS



static

dynamic



Static



Small overhead



Space efficient



Can unload modules



Complex



Bugs

Dynamic

Simple implementation

Can support more platforms

Low maintenance

Loading is costly

Accumulates loaded modules

Recap



Hack on GHC

Help fix those bugs^b

Fun and rewarding it is!

^b Looking at you [GHC-13624](#).

Thanks for your attention!

Questions?

- speaker: Artem Pianykh
- twitter: @artem_pyanykh

